

Harshal V. Abhyankar

✉ abhyankarharshal22@gmail.com | ☎ 217-926-3518 | 📠 harshal-a-omni | 🌐 omniwot | 📍 Charlotte, NC

SUMMARY: Forward-deployed AI engineer shipping production agentic systems into regulated enterprise environments — RAG, LLMOps, guardrails, and evaluation frameworks that gate releases on accuracy, latency, and cost.

EDUCATION

Information Management — <i>Master of Science</i> University of Illinois at Urbana-Champaign	AUG 2023 - MAY 2025 GPA: 4.00/4.00
Electrical and Electronics — <i>Bachelor of Engineering</i> Birla Institute of Technology and Science, Pilani	AUG 2019 - MAY 2023 GPA: 3.59/4.00

TECHNICAL SKILLS

Languages: Python, SQL, TypeScript, JavaScript, R, C++
GenAI & Agentic: RAG, LangGraph, LangChain, Mastra, MCP, Prompt Engineering, Function Calling / Tool Use, PEFT / LoRA, PyTorch
Evals & LLMOps: Langfuse, Ragas, DeepEval, LLM-as-Judge, OpenTelemetry
Cloud & Platform: Azure, AWS (Bedrock, Lambda, S3), GCP, Kubernetes / OpenShift, Docker, CI/CD, FastAPI, REST APIs
Data: PostgreSQL, OracleDB, Weaviate, Pinecone, ArangoDB, Elasticsearch
Voice & Realtime: ElevenLabs, Twilio, OpenAI Realtime API
Certifications: Microsoft Azure Solutions Architect Expert (AZ-305), Deep Learning Specialization (Coursera)

WORK EXPERIENCE

Theorem Labs <i>Programmer Analyst, Forward Deployed Engineering</i>	AUG 2025 - PRESENT <i>Charlotte, NC</i>
--	--

- Lead a 3-person forward-deployed pod delivering AI capabilities into MetricStream’s governance, risk, and compliance (GRC) platform; converted a pre-sales PoC into a signed \$75K pilot and a \$600K production deployment (\$675K TCV).
- Embed prompt guardrails, audit logging, and data-handling controls to meet enterprise governance and auditability requirements; own pilot and production delivery end-to-end from requirements discovery through deployment, turning around customer pilots and demos on 2–3 week build cycles.
- Instrument Langfuse tracing and custom evaluation pipelines that gate every production release on hallucination rate, citation grounding, latency, and cost per run, catching regressions before enterprise handoff.
- Architect full-stack AI applications on Python, AWS, Mastra, and React; build reusable AI development harnesses (Cursor rules, Claude Code skills) adopted team-wide to standardize LLM-assisted delivery.

Very Sweet Inc. <i>AI Engineer</i>	DEC 2024 - MAR 2025 <i>Remote</i>
--	--------------------------------------

- Engineered agentic workflows combining GCP Document AI OCR with LLM-driven parsing to automate large-scale document intake, reducing manual processing effort by 90%.
- Partnered with front-end engineers to ship AI document-processing workflows into production; designed the OCR-to-structured-data API contract enabling parallel frontend and backend development.

BayInfotech <i>AI Software Engineer</i>	MAY 2024 - JUL 2024 <i>Livermore, CA</i>
---	---

- Optimized RAG pipelines by benchmarking retrieval models and tuning vector search across Elasticsearch, LangChain, Ray.io, and Pinecone.
- Built distributed data ingestion and processing pipelines for multi-million record datasets using GCP and Ray.io, reducing ETL runtimes by 40%.

TECHNICAL PROJECTS

Document Analysis Agent for Airline Pilots — Business Intelligence Group, UIUC	JAN 2025 - MAY 2025
--	---------------------

- Led a team to ship a production AI agent for querying Collective Bargaining Agreements in natural language, combining RAG, graph-based search, and LLM reasoning; built an LLM-as-judge eval harness to validate citation accuracy before deployment.
- Translated legal-domain stakeholder requirements into system prompt guardrails and retrieval filters constraining outputs to verified contract clauses, preventing hallucinated contractual claims.