

Harshal V. Abhyankar

✉ abhyankarharshal22@gmail.com | ☎ 217-926-3518 | 🌐 harshal-a-omni | 🌐 omniwot | 📍 Charlotte, NC

SUMMARY: Forward-deployed AI engineer building agentic systems and enterprise AI prototypes — deep in Mastra, LLMOps, and prompt engineering; shipping from whiteboard to production.

EDUCATION

Information Management — *Master of Science* AUG 2023 - MAY 2025
University of Illinois at Urbana-Champaign GPA: 4.00/4.00

Electrical and Electronics — *Bachelor of Engineering* AUG 2019 - MAY 2023
Birla Institute of Technology and Science, Pilani GPA: 3.59/4.00

TECHNICAL SKILLS

Languages: Python, SQL, R, JavaScript, TypeScript, C, C++, LCEL

Agentic & AI: Mastra, LangGraph, LangChain, MCP, RAG Agents, Prompt Engineering, PEFT / LoRA, PyTorch

Voice & Realtime: ElevenLabs, Twilio, OpenAI Realtime API

Cloud, Infra & Data: AWS (Lambda, S3, Bedrock), GCP, Azure, ArangoDB, Weaviate CE, Pinecone, PostgreSQL, OracleDB, FastAPI, Docker, Kubernetes, Node.js, React.js

LLMOps, Evals & Tooling: Langfuse, Ragas, DeepEval, LLM-as-Judge, Lovable, Cursor, Claude Code

Certifications: Microsoft Azure Solutions Architect Expert (AZ-305), Deep Learning Specialization (Coursera)

WORK EXPERIENCE

Theorem Labs AUG 2025 - PRESENT
Programmer Analyst, Forward Deployed Engineering Charlotte, NC

- Lead a 3-person forward-deployed pod for MetricStream, driving AI PoC partnerships and GRC pilot contracts valued at over \$500K in total contract value.
- Own end-to-end prototype delivery — from requirements discovery through deployment — embedding prompt guardrails, audit logging, and data handling controls to meet enterprise governance requirements.
- Architect full-stack AI applications using Python, AWS, Mastra, and React; instrument Langfuse tracing and custom eval pipelines to gate releases on latency, hallucination rate, and cost per run before enterprise handoff.
- Build and maintain AI development harnesses — Cursor rules and Claude Code skills — adopted across the team to standardize LLM-assisted workflows and accelerate delivery on every project.

Very Sweet Inc. DEC 2024 - MAR 2025
AI Engineer Remote

- Engineered agentic workflows combining GCP Document AI OCR with LLM-driven parsing to automate large-scale document intake, reducing manual processing effort by 90%.
- Partnered with front-end engineers to ship AI document-processing workflows into production; designed the OCR-to-structured-data API contract enabling parallel frontend and backend development, accelerating prototype-to-launch timelines.

BayInfotech MAY 2024 - JUL 2024
AI Software Engineer Livermore, CA

- Optimized RAG pipelines by benchmarking retrieval models and tuning vector search across Elasticsearch, LangChain, Ray.io, and Pinecone.
- Built distributed data ingestion and processing pipelines for multi-million record datasets using GCP and Ray.io, reducing ETL runtimes by 40%.

TECHNICAL PROJECTS

Promptline — Conversational Voice AI Platform ([promptline.app](#)) 2024 - 2025
• Architected and built Promptline end-to-end — a voice AI platform for customer service over phone and text using ElevenLabs, Twilio, and Lovable; adopted by TheoremLabs as a core product accelerator, now maintained by a dedicated team.

Document Analysis Agent for Airline Pilots — Project Manager, Business Intelligence Group, UIUC JAN 2025 - MAY 2025
• Led a team to ship a production AI agent for querying Collective Bargaining Agreements in natural language, combining RAG, graph-based search, and LLM reasoning; built an LLM-as-judge eval harness to validate citation accuracy before deployment.