

Harshal V. Abhyankar

✉ abhyankarharshal22@gmail.com | ☎ 217-926-3518 | 🌐 harshal-a-omni | 🌀 omniwot | 📍 Charlotte, NC

SUMMARY: Forward-deployed AI engineer building agentic systems and enterprise AI prototypes — deep in Mastra, LLMOps, and prompt engineering; shipping from whiteboard to production.

EDUCATION

Information Management — <i>Master of Science</i> University of Illinois at Urbana-Champaign	AUG 2023 - MAY 2025 GPA: 4.00/4.00
Electrical and Electronics — <i>Bachelor of Engineering</i> Birla Institute of Technology and Science, Pilani	AUG 2019 - MAY 2023 GPA: 3.59/4.00

TECHNICAL SKILLS

Languages: Python, SQL, R, JavaScript, TypeScript, C, C++, LCEL
Agentic & AI: Mastra, LangGraph, LangChain, MCP, RAG Agents, Prompt Engineering, PEFT / LoRA, PyTorch
Voice & Realtime: ElevenLabs, Twilio, OpenAI Realtime API
Cloud, Infra & Data: AWS (Lambda, S3, Bedrock), GCP, Azure, ArangoDB, Weaviate CE, Pinecone, PostgreSQL, OracleDB, FastAPI, Docker, Kubernetes, Node.js, React.js
LLMOps, Evals & Tooling: Langfuse, Ragas, DeepEval, LLM-as-Judge, Lovable, Cursor, Claude Code
Certifications: Microsoft Azure Solutions Architect Expert (AZ-305), Deep Learning Specialization (Coursera)

WORK EXPERIENCE

Theorem Labs <i>Programmer Analyst, Forward Deployed Engineering</i>	AUG 2025 - PRESENT Charlotte, NC
- Lead a 3-person forward-deployed pod for MetricStream; converted a pre-sales PoC into a signed \$75K pilot, then delivered it successfully into a \$600K production deployment (\$675K TCV). - Serve as lead developer and solution architect, owning pilot and production delivery end-to-end — from requirements discovery through deployment — embedding prompt guardrails, audit logging, and data handling controls to meet enterprise governance requirements. - Architect full-stack AI applications using Python, AWS, Mastra, and React; instrument Langfuse tracing and custom eval pipelines to gate releases on latency, hallucination rate, and cost per run before enterprise handoff. - Build and maintain AI development harnesses — Cursor rules and Claude Code skills — adopted across the team to standardize LLM-assisted workflows and accelerate delivery on every project.	
Very Sweet Inc. <i>AI Engineer</i>	DEC 2024 - MAR 2025 Remote
- Engineered agentic workflows combining GCP Document AI OCR with LLM-driven parsing to automate large-scale document intake, reducing manual processing effort by 90%. - Partnered with front-end engineers to ship AI document-processing workflows into production; designed the OCR-to-structured-data API contract enabling parallel frontend and backend development, accelerating prototype-to-launch timelines.	
BayInfotech <i>AI Software Engineer</i>	MAY 2024 - JUL 2024 Livermore, CA
- Optimized RAG pipelines by benchmarking retrieval models and tuning vector search across Elasticsearch, LangChain, Ray.io, and Pinecone. - Built distributed data ingestion and processing pipelines for multi-million record datasets using GCP and Ray.io, reducing ETL runtimes by 40%.	

TECHNICAL PROJECTS

Promptline — Conversational Voice AI Platform (promptline.app)	2024 - 2025
Architected and built Promptline end-to-end — a voice AI platform for customer service over phone and text using ElevenLabs, Twilio, and Lovable; adopted by TheoremLabs as a core product accelerator, now maintained by a dedicated team.	
Document Analysis Agent for Airline Pilots — Project Manager, Business Intelligence Group, UIUC	JAN 2025 - MAY 2025
Led a team to ship a production AI agent for querying Collective Bargaining Agreements in natural language, combining RAG, graph-based search, and LLM reasoning; built an LLM-as-judge eval harness to validate citation accuracy before deployment.	